

# A (Not So) Simple (Full) Matrix Class

April 5, 2016

# Exercise

Reorganize the `fem1d` exercise code in the following way

- ▶ Allow the user to specify the problem and algorithm parameters on the command line rather than in a header
  - ▶ Use `GetPot`
- ▶ Refactor the code by implementing a separate library for matrices (and column vectors implemented as 1-column matrices) based on STL containers such as `std::vector<>` or `std::array<>` with the following methods/functions
  - ▶ `transpose` :  $A = A^T$
  - ▶ `gauss-seidel` : solve  $Ax = b$  iteratively as in previous versions of the example
  - ▶ `solve` : solve  $Ax = b$  by means of a direct method
  - ▶ `operator*` : matrix-matrix and matrix-vector multiplication

# Solution I

- ▶ `matrix-0.1/matrix.{h,cpp}`
  - ▶ `std::array<>` requires the size to be known at compile time, in order to make the matrix object configurable at run time we use `std::vector<>`
  - ▶ matrices are organized as *column major*, i.e.  $A(i, j) = \text{data}[i + j * \text{rows}()]$ , conversion from 1d to 2d indexing is performed by the utility (private) function `sub2ind`
  - ▶ access to elements is implemented both in const and non-const versions, by overloading **operator()**
  - ▶ data is private, *getter methods* expose what is needed to the user, both const and non-const versions are provided
  - ▶ prototypes of fortran77 functions defined as **extern "C" {...}**, assume the compiler adds an underscore (this may change depending on the compiler), define upper and lower case versions

# Solution II

- ▶ use precompiler conditionals to let user select at compile time which implementation to use
- ▶ naive implementation of matrix-matrix multiplication is slow because it has low *data locality* (see below more about memory hierarchy and cache optimized algorithms), simply transposing the left matrix factor improves performance significantly
- ▶ BLAS and LAPACK provide cache optimized linear algebra subroutines, optimizations are hardware dependent, the interface of BLAS and LAPACK is fixed, different implementations with different performance are available
  - ▶ ATLAS does self tuning via experiments
  - ▶ OpenBlas has hand-tuned versions for most common architectures
  - ▶ vendor-specific implementations (Intel MKL, Apple vecLib, ...)
- ▶ **#include** <ctime> header provides timing utilities, `tic ()` and `toc (x)` macros start and stop the timer (like in Matlab)

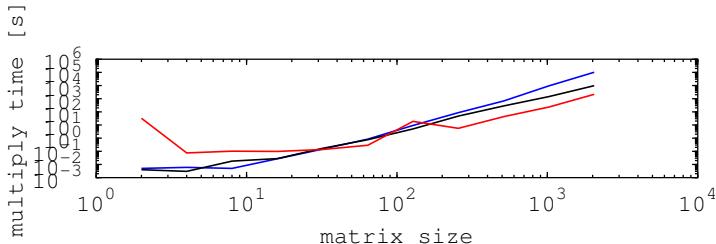
▶ `matrix-0.1/Makefile`

# Solution III

- ▶ build static and/or shared library
- ▶ user provided flags are passed on to the preprocessor, compiler and linker
- ▶ `matrix-0.1/test_*.cpp`
  - ▶ for each feature a test to verify consistency and/or performance is provided
  - ▶ `run_test.m` builds `test_matrix_mult` with various options and compares performance of various implementations
- ▶ `fem1d-0.4/fem1d.h`
  - ▶ define help text using `std::string`
  - ▶ remove `config.h`
- ▶ `fem1d-0.4/fem1d.cpp`
  - ▶ define problem/algorithm constants using GetPot methods

# Solution IV

- ▶ use matrix objects to define matrices and vectors
- ▶ `fem1d-0.4/Makefile`
  - ▶ include flags for matrix headers and libraries



**Figure:** Performance of the different implementations of the matrix multiplication: 1) standard (blue), 2) with transpose of first factor (black), 3) DGEMM (red). For a matrix of size  $2^{11}$  the algorithm 2) is 10 times faster than 1) and algorithm 3) is 64 times faster. The comparison was run on MacBook Pro with core2-duo CPU, code was compiled with g++ 4.8 and linked to OpenBlas.

# Matrix Class (Header)

Lines 1–25 / 85

```
1  #ifndef HAVE_MATRIX_H
2  #define HAVE_MATRIX_H
3
4  #include <vector>
5  #include <iostream>
6
7  class
8  matrix
9  {
10
11  private :
12
13      std::vector<double> data;
14      const unsigned int rows;
15      const unsigned int cols;
16
17      inline
18      unsigned int
19      sub2ind (const unsigned int ir ,
20              const unsigned int jc) const
21      { return (ir + jc * rows); };
22
23      double &
24      index (unsigned int irow, unsigned int jcol)
25      { return data[sub2ind (irow, jcol)]; };
```



# Matrix Class (Header)

Lines 25–49 / 85

```
25     { return data[sub2ind (irow, jcol)]; };
26
27     const double &
28     const_index (unsigned int irow, unsigned int jcol) const
29     { return data[sub2ind (irow, jcol)]; };
30
31     public :
32
33     matrix (unsigned int size)
34         : rows (size), cols (size)
35     { data.resize (rows * cols, 0.0); };
36
37     matrix (unsigned int rows_,
38             unsigned int cols_)
39         : rows (rows_), cols (cols_)
40     { data.resize (rows * cols, 0.0); };
41
42     matrix (matrix const &) = default;
43
44     unsigned int
45     get_rows () const { return rows; }
46
47     unsigned int
48     get_cols () const { return cols; }
```

# Matrix Class (Header)

Lines 49–73 / 85

```
49
50 double &
51 operator() (unsigned int irow, unsigned int jcol)
52 { return index (irow, jcol); };
53
54 const double &
55 operator() (unsigned int irow, unsigned int jcol) const
56 { return const_index (irow, jcol); };
57
58 const double *
59 get_data () const { return &(data[0]); };
60
61 double *
62 get_data () { return &(data[0]); };
63
64 /// transposed matrix : B = A'
65 matrix
66 transpose () const;
67
68 void
69 solve (matrix &rhs);
70
71 void
72 gauss_seidel
73 (const matrix &f,
```

# Matrix Class (Header)

Lines 73–97 / 85

```
73     (const matrix &f,  
74      matrix &uh,  
75      const unsigned int maxit = 100,  
76      const double tol = 1.e-9);  
77  
78 };  
79  
80 /// matrix x matrix product : C = A * B  
81 matrix  
82 operator* (const matrix& A, const matrix& B);  
83  
84 #endif
```

# Matrix Class (Implementation)

Lines 1–25 / 139

```
1  #include "matrix.h"
2  #include <cassert>
3  #include <algorithm>
4  #include <ctime>
5
6  #define dgesv dgesv_
7  #define DGESV dgesv_
8
9  #define dgemm dgemm_
10 #define DGEMM dgemm_
11
12 extern "C"
13 {
14     void
15     dgesv (const int *N, const int *NRHS, const double *A,
16           const int *LDA, int *IPIV, double *B, const int *LDB,
17           int *INFO);
18
19     void
20     dgemm (const char *TRANSA, const char *TRANSB, const int *M,
21           const int *N, const int *K, const double *ALPHA,
22           const double *A, const int *LDA, const double *B,
23           const int *LDB, const double *BETA, double *C,
24           const int *LDC);
25 }
```

# Matrix Class (Implementation)

Lines 25–49 / 139

```
25 }
26
27
28 matrix
29 matrix::transpose () const
30 {
31     matrix retval (get_cols (), get_rows ());
32     unsigned int i, j;
33     for (i = 0; i < rows; ++i)
34         for (j = 0; j < cols; ++j)
35             retval (i, j) = const_index (j, i);
36     return (retval);
37 }
38
39 #ifndef USE_DGEMM
40 matrix
41 operator* (const matrix& A, const matrix& B)
42 {
43
44     int M = A.get_rows ();
45     int N = B.get_cols ();
46     int K = A.get_cols ();
47     assert (K == B.get_rows ());
48
49     char ntr = 'n';
```

# Matrix Class (Implementation)

Lines 49–73 / 139

```
49     char ntr = 'n';
50     double one = 1.0;
51     double zero = 0.0;
52
53     matrix retval (M, N);
54
55     dgemm (&ntr, &ntr, &M, &N, &K,
56           &one, A.get_data (), &M,
57           B.get_data (), &K, &zero,
58           retval.get_data (), &M);
59
60     return (retval);
61 }
62 #elif defined (MAKE_TMP_TRANSPOSE)
63 matrix
64 operator* (const matrix& A, const matrix& B)
65 {
66     assert (A.get_cols () == B.get_rows ());
67     matrix retval (A.get_rows (), B.get_cols ());
68     matrix tmp = A.transpose ();
69     for (unsigned int i = 0; i < retval.get_rows (); ++i)
70         for (unsigned int j = 0; j < retval.get_cols (); ++j)
71             for (unsigned int k = 0; k < A.get_cols (); ++k)
72                 retval(i,j) += tmp(k,i) * B(k,j);
73     return (retval);
```

# Matrix Class (Implementation)

Lines 73–97 / 139

```
73     return (retval);
74 }
75 #else
76 matrix
77 operator* (const matrix& A, const matrix& B)
78 {
79     assert (A.get_cols () == B.get_rows ());
80     matrix retval (A.get_rows (), B.get_cols ());
81     for (unsigned int i = 0; i < retval.get_rows (); ++i)
82         for (unsigned int j = 0; j < retval.get_cols (); ++j)
83             for (unsigned int k = 0; k < A.get_cols (); ++k)
84                 retval(i,j) += A(i,k) * B(k,j);
85     return (retval);
86 }
87 #endif
88
89 void
90 matrix::solve (matrix &rhs)
91 {
92     int N = get_rows ();
93     int IPIV[N];
94     int NRHS = rhs.get_cols ();
95     int LDB = rhs.get_rows ();
96     int INFO = 0;
97     matrix tmp ((*this));
```

# Matrix Class (Implementation)

Lines 97–121 / 139

```
97     matrix tmp ((*this));
98     dgesv (&N, &NRHS, tmp.get_data (), &N,
99           IPIV, rhs.get_data (), &LDB, &INFO);
100 };
101
102
103 void
104 matrix::gauss_seidel
105 (const matrix &f,
106  matrix &uh,
107   const unsigned int maxit,
108   const double tol)
109 {
110
111     assert (uh.get_rows () == f.get_rows ());
112     assert (uh.get_rows () == get_rows ());
113     assert (uh.get_cols () == 1);
114
115     double uh_new = 0;
116     double incrnorm = 0;
117
118     for (unsigned int ii = 0; ii < maxit; ++ii)
119     {
120         incrnorm = 0;
121         for (unsigned int jj = 0; jj < get_rows (); ++jj)
```



# Matrix Class (Implementation)

Lines 121–145 / 139

```
121     for (unsigned int jj = 0; jj < get_rows (); ++jj)
122     {
123         double res = f(jj , 0);
124
125         for (unsigned int kk = 0; kk < get_cols (); ++kk)
126             if (kk != jj)
127                 res -= const_index(jj ,kk) * uh(kk, 0);
128
129         uh_new = res / const_index(jj ,jj);
130         incrnorm = std::abs (uh_new - uh(jj , 0)) > incrnorm ?
131             std::abs (uh_new - uh(jj , 0)) :
132             incrnorm;
133         uh(jj , 0) = uh_new;
134     }
135     if (incrnorm < tol)
136         break;
137 }
138 };
```

# Matrix Class (Test Multiply)

Lines 1–24 / 24

```
1  #include "matrix.h"
2  #include "test_matrix_mult.h"
3  int main ()
4  {
5
6      matrix A(msize);
7      matrix B(msize);
8      for (unsigned int i = 0; i < msize; ++i)
9      {
10         A(i,i) = 10.0;
11         A(i,msize-1) = 30.0;
12         B(i,0) = 1.0;
13         B(i,i) = 3.0;
14     }
15
16     std::cout << "msize_=" << msize << std::endl;
17
18     tic ();
19     matrix C = A * B;
20     toc ("multiply_time_=");
21
22     return 0;
23 }
```

# Matrix Class (Test Gauß-Seidel)

Lines 1–25 / 37

```
1  #include "test_gauss_seidel.h"
2
3  int main ()
4  {
5
6      matrix A (10);
7      matrix rhs (10, 1);
8
9      for (unsigned int i = 0; i < 10; ++i)
10     {
11         rhs (i, 0) = 1.;
12         for (unsigned int j = 0; j < 10; ++j)
13         {
14             if (j+1 == i)
15                 A(i, j) = -1.;
16             else if (j == i+1)
17                 A(i, j) = -1.;
18             else if (i == j)
19                 A(i, j) = 2.;
20         }
21     }
22
23     matrix uh (rhs);
24     A.gauss_seidel (rhs, uh, 1000);
```

# Matrix Class (Test Gauß-Seidel)

Lines 25–49 / 37

```
25
26     matrix f = A * uh;
27
28     for (unsigned int i = 0; i < 10; ++i)
29     {
30         std::cout << "x_(" << i+1 << ") = " << uh(i, 0) << "; "
31                 << "y_(" << i+1 << ") = " << f(i, 0) << "; "
32                 << "b_(" << i+1 << ") = " << rhs(i, 0) << "; "
33                 << std::endl;
34     }
35     return 0;
36 }
```

# Matrix Class (Test Solve)

Lines 1–25 / 38

```
1  #include "test_solve.h"
2  #include <iostream>
3
4  int main ()
5  {
6
7      matrix A (10);
8      matrix rhs (10, 1);
9
10     for (unsigned int i = 0; i < 10; ++i)
11     {
12         rhs (i, 0) = 1.;
13         for (unsigned int j = 0; j < 10; ++j)
14         {
15             if (j+1 == i)
16                 A(i, j) = -1.;
17             else if (j == i+1)
18                 A(i, j) = -1.;
19             else if (i == j)
20                 A(i, j) = 2.;
21         }
22     }
23
24     matrix uh (rhs);
25     A.solve (uh);
```

# Matrix Class (Test Solve)

Lines 25–49 / 38

```
25     A.solve (uh);
26
27     matrix f = A * uh;
28
29     for (unsigned int i = 0; i < 10; ++i)
30     {
31         std::cout << "x_(" << i+1 << ") = " << uh(i, 0) << ";"
32                 << "      y_(" << i+1 << ") = " << f(i, 0) << ";"
33                 << "      b_(" << i+1 << ") = " << rhs(i, 0) << ";"
34                 << std::endl;
35     }
36     return 0;
37 }
```

# Matrix Class (run\_test.m)

Lines 1–25 / 57

```
1
2
3  fmt = ...
4  'make_CXX="g++" _CPPFLAGS="-I. -Usize -Dmsize=%d_%s" _distclean_test_matrix_mult\'
5
6  m  = [];
7  mm = [];
8  mmm = [];
9  d  = [];
10
11  for ii = 1 : 11
12
13      printf (fmt, 2^ii, "");
14      [a, str] = system (sprintf (fmt, 2^ii, ""));
15      [a, str] = system ('./test_matrix_mult');
16      if (a != 0)
17          disp (str)
18          return
19      endif
20      eval (str);
21      d = [d; msize];
22      m = [m; multiply_time];
23
24      printf (fmt, 2^ii, "-DMAKE_TMP TRANSP");
25      [a, str] = system (sprintf (fmt, 2^ii, "-DMAKE_TMP TRANSP"));
```

# Matrix Class (run\_test.m)

Lines 25–49 / 57

```
25 [a, str] = system (sprintf (fmt, 2^ii, "-DMAKE_TMP_TRANSP"));
26 [a, str] = system ('./test_matrix_mult');
27 if (a != 0)
28     disp (str)
29     return
30 endif
31 eval (str);
32 mm = [mm; multiply_time];
33
34 printf (fmt, 2^ii, "-DUSE_DGEMM");
35 [a, str] = system (sprintf (fmt, 2^ii, "-DUSE_DGEMM"));
36 [a, str] = system ('./test_matrix_mult');
37 if (a != 0)
38     disp (str)
39     return
40 endif
41 eval (str);
42 mmm = [mmm; multiply_time];
43
44 loglog (d, m, 'linewidth', 5,
45         d, mm, 'linewidth', 5,
46         d, mmm, 'linewidth', 5);
47
48 set (gca (), "FontSize", 18)
49 legend ("without_transpose", "with_transpose", "with_dgemm",
```



# Matrix Class (run\_test.m)

Lines 49–73 / 57

```
49     legend ("without_transpose", "with_transpose", "with_dgemm",  
50            "location", "southoutside")  
51     xlabel ("matrix_size", "FontSize", 18)  
52     ylabel ("multiply_time_[ms]", "FontSize", 18)  
53     axis tight  
54     drawnow  
55     endfor  
56  
57     print -dpdf speedup.pdf
```

# Matrix Class Library Makefile

Lines 1–25 / 43

```
1  OBJS = matrix.o
2  TEST_OBJS = test_matrix_mult.o test_gauss_seidel.o test_solve.o
3  TESTS = test_matrix_mult test_gauss_seidel test_solve
4  HEADERS = matrix.h
5  SHARED_EXT = so
6
7  .PHONY : tests
8
9  tests : $(TESTS)
10
11 libmatrix.a : $(OBJS)
12     $(AR) crv libmatrix.a $(OBJS)
13
14 libmatrix.$(SHARED_EXT) : $(OBJS)
15     $(CXX) -shared -L. $(LDFLAGS) -o $$@ $(OBJS) $(LIBS)
16
17 test_matrix_mult : test_matrix_mult.o libmatrix.a
18     $(CXX) -L. $(LDFLAGS) -o $$@ $< -lmatrix $(LIBS)
19
20 test_gauss_seidel : test_gauss_seidel.o libmatrix.a
21     $(CXX) -L. $(LDFLAGS) -o $$@ $< -lmatrix $(LIBS)
22
23 test_solve : test_solve.o libmatrix.a
24     $(CXX) -L. $(LDFLAGS) -o $$@ $< -lmatrix $(LIBS)
```

# Matrix Class Library Makefile

Lines 25–49 / 43

```
25
26 speedup.pdf : run_test.m
27             octave $<
28
29 $(OBSJ) : %.o : %.cpp %.h
30             $(CXX) $(CPPFLAGS) $(CXXFLAGS) -c $<
31
32 $(TEST_OBSJ) : %.o : %.cpp %.h
33             $(CXX) $(CPPFLAGS) $(CXXFLAGS) -c $<
34
35 .PHONY: clean distclean
36
37 clean :
38         $(RM) $(OBSJ) $(TEST_OBSJ)
39
40 distclean : clean
41         $(RM) $(TESTS) \
42         speedup.pdf libmatrix.a libmatrix.$(SHARED.EXT)
```

# fem1d main

Lines 1–25 / 79

```
1  #include "fem1d.h"
2  #include "matrix.h"
3  #include "GetPot"
4
5  int main (int argc, char **argv)
6  {
7
8      GetPot cl (argc, argv);
9
10     if (cl.search (2, "-h", "--help"))
11     {
12         std::cerr << help_text << std::endl;
13         return 0;
14     }
15     const double a = cl.follow (0.0, "-a");
16     const double b = cl.follow (1.0, "-b");
17     const unsigned int nnodes = cl.follow (100, 2, "-n", "-nnodes");
18
19     mesh m (a, b, nnodes);
20     matrix A(nnodes);
21
22     matrix mloc(2);
23     for (unsigned int iel = 0; iel < m.nels; ++iel)
24     {
```

# fem1d main

Lines 25–49 / 79

```
25
26         std::fill (mloc.get_data (),
27                     mloc.get_data () + 4,
28                     0.0);
29
30     for (unsigned int inode = 0; inode < 2; ++inode)
31     {
32         double igrad = (inode == 0 ? 1.0 / m.h : -1.0 / m.h);
33         for (unsigned int jnode = 0; jnode < 2; ++jnode)
34         {
35             double jgrad = (jnode == 0 ? 1.0 / m.h : -1.0 / m.h);
36             mloc(inode, jnode) = igrad * jgrad * m.h;
37             A(m.elements[iel][inode], m.elements[iel][jnode]) +=
38                 mloc(inode, jnode);
39         }
40     }
41 }
42
43 matrix f(nnodes, 1);
44 matrix vloc(2, 1);
45
46 for (unsigned int iel = 0; iel < m.nels; ++iel)
47 {
48     std::fill (mloc.get_data (),
49                 mloc.get_data () + 2,
```

# fem1d main

Lines 49–73 / 79

```
49         mloc.get_data () + 2,  
50         0.0);  
51  
52     for (unsigned int inode = 0; inode < 2; ++inode)  
53     {  
54         vloc(inode, 0) = m.h / 2.0;  
55         f(m.elements[iel][inode], 0) += vloc(inode, 0);  
56     }  
57 }  
58  
59 f(0, 0) = 0;  
60 f(nnodes - 1, 0) = 0;  
61  
62 A(0,0) = 1.0;  
63 A(nnodes-1,nnodes-1) = 1.0;  
64 for (unsigned int ii = 1; ii < nnodes; ++ii)  
65 {  
66     A(0, ii) = 0.0;  
67     A(nnodes-1, nnodes-1-ii) = 0.0;  
68 }  
69  
70  
71 matrix uh(f);  
72 A.gauss_seidel (f, uh);
```

# fem1d main

Lines 73–97 / 79

```
73
74     for (unsigned int ii = 0; ii < nnodes; ++ii)
75         std::cout << uh(ii, 0) << std::endl;
76
77     return 0;
78     };
```

# fem1d Header

Lines 1–24 / 24

```
1  #ifndef HAVE_FEM1D_H
2  #define HAVE_FEM1D_H
3
4  #include <array>
5  #include <algorithm>
6  #include <cmath>
7  #include <iostream>
8  #include <string>
9
10 #include "mesh.h"
11
12
13 std::string help_text = std::string () +
14 "fem1d_command_line_options\n" +
15   "-h, --help          print this text and exit\n" +
16   "-a <value>          first end of interval\n" +
17   "-b <value>          second end of interval\n" +
18   "-n, --nnodes <value> number of triangulation nodes\n" +
19   "-m, --maxit <value> number of iterations\n" +
20   "-t, --tol <value>   tolerance\n";
21
22
23 #endif
```



# fem1d Executable Makefile

Lines 1–20 / 20

```
1  OBJS = fem1d.o mesh.o
2  HEADERS = config.h
3  CPPFLAGS ?= -I../matrix -0.1/ -I.
4  LDFLAGS ?= -L../matrix -0.1/
5  LIBS =
6
7  fem1d : $(OBJS)
8          $(CXX) $(LDFLAGS) $(OBJS) -o $$@ -lmatrix $(LIBS)
9
10 $(OBJS) : %.o : %.cpp %.h $(HEADERS)
11          $(CXX) $(CPPFLAGS) $(CXXFLAGS) -c $<
12
13 .PHONY: clean distclean
14
15 clean :
16          $(RM) $(OBJS)
17
18 distclean : clean
19          $(RM) fem1d
```

# Memory hierarchy and Cache Optimized Algorithms

# Memory hierarchy and Cache Optimized Algorithms

insert pages 1-9 of file images/slides.pdf

# BLAS and LAPACK

# BLAS and LAPACK

insert pages 10-20 of file images/slides.pdf